



```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

```
In [2]: Boston_df = pd.read_csv("HousingData.csv")
```

```
In [3]: print(Boston_df.head())
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```
In [4]: print(Boston_df.shape)
```

(506, 14)

Data preprocessing

```
In [5]: Boston_df.isnull().sum()
```

```
Out[5]: CRIM      20
ZN          20
INDUS       20
CHAS        20
NOX         0
RM          0
AGE         20
DIS         0
RAD         0
TAX         0
PTRATIO     0
B           0
LSTAT       20
MEDV        0
dtype: int64
```

```
In [6]: fig, ax = plt.subplots(figsize=(8,8))
sns.distplot(Boston_df.CRIM)
```

```
plt.show()

fig, ax = plt.subplots(figsize=(8,8))
sns.distplot(Boston_df.ZN)
plt.show()

fig, ax = plt.subplots(figsize=(8,8))
sns.distplot(Boston_df.INDUS)
plt.show()

fig, ax = plt.subplots(figsize=(8,8))
sns.distplot(Boston_df.CHAS)
plt.show()

fig, ax = plt.subplots(figsize=(8,8))
sns.distplot(Boston_df.AGE)
plt.show()

fig, ax = plt.subplots(figsize=(8,8))
sns.distplot(Boston_df.LSTAT)
plt.show()
```


C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\1919813569.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Boston_df.CRIM)
```




C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\1919813569.py:6: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Boston_df.CRIM)
```




C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\1919813569.py:10: UserWarning:

``distplot` is a deprecated function and will be removed in seaborn v0.14.0.`

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Boston_df.INDUS)
```




C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\1919813569.py:14: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Boston_df.CHAS)
```




C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\1919813569.py:18: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Boston_df.AGE)
```




C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\1919813569.py:22: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(Boston_df.LSTAT)
```



```
In [7]: # handling the missing values
# Impute missing values in the 'CRIM' column with the median
median_crim = Boston_df['CRIM'].median()
Boston_df['CRIM'].fillna(median_crim, inplace=True)

# Impute missing values in the 'ZN' column with the median
median_zn = Boston_df['ZN'].median()
Boston_df['ZN'].fillna(median_zn, inplace=True)

# Impute missing values in the 'INDUS' column with the median
median_indus = Boston_df['INDUS'].median()
Boston_df['INDUS'].fillna(median_indus, inplace=True)

# using the mode for imputation is a suitable approach. This will replace miss
# Impute missing values in the 'CHAS' column with the mode
mode_chas = Boston_df['CHAS'].mode()[0] # .mode() can return multiple values i
Boston_df['CHAS'].fillna(mode_chas, inplace=True)
```



```
# Impute missing values in the 'AGE' column with the median
median_age = Boston_df['AGE'].median()
Boston_df['AGE'].fillna(median_age, inplace=True)

# Impute missing values in the 'LSTAT' column with the median
median_lstat = Boston_df['LSTAT'].median()
Boston_df['LSTAT'].fillna(median_lstat, inplace=True)
```

C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\2845753325.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
Boston_df['CRIM'].fillna(median_crim, inplace=True)
```

C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\2845753325.py:8: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
Boston_df['ZN'].fillna(median_zn, inplace=True)
```

C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\2845753325.py:12: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
Boston_df['INDUS'].fillna(median_indus, inplace=True)
```

C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\2845753325.py:17: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
Boston_df['CHAS'].fillna(mode_chas, inplace=True)
```

C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\2845753325.py:21: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

ned assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
Boston_df['AGE'].fillna(median_age, inplace=True)
```

C:\Users\erraj\AppData\Local\Temp\ipykernel_38320\2845753325.py:25: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
Boston_df['LSTAT'].fillna(median_lstat, inplace=True)
```

```
In [8]: Boston_df.isnull().sum()
```

```
Out[8]: CRIM      0
        ZN        0
        INDUS    0
        CHAS     0
        NOX      0
        RM       0
        AGE      0
        DIS      0
        RAD      0
        TAX      0
        PTRATIO  0
        B        0
        LSTAT    0
        MEDV     0
        dtype: int64
```

```
In [9]: # all the statistical method at once
        Boston_df.describe()
```

Out[9]:


	CRIM	ZN	INDUS	CHAS	NOX	RM
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.479140	10.768775	11.028893	0.067194	0.554695	6.284634
std	8.570832	23.025124	6.704679	0.250605	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.083235	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.253715	0.000000	9.690000	0.000000	0.538000	6.208500
75%	2.808720	0.000000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

Cheacking of possitive and Negative correlation

```
In [10]: correlation = Boston_df.corr()
```

```
In [11]: # constructing heatmap to understand the correlation
plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_
```

```
Out[11]: <Axes: >
```



splitting the data and Target

```
In [12]: X = Boston_df.drop(['MEDV'], axis=1)    # dropping column so axis = 1
         Y = Boston_df['MEDV']
```

```
In [13]: print(X)
         print(Y)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0.0	0.573	6.030	76.8	2.5050	1	273	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	11.43
..
501	21.0	391.99	11.43
502	21.0	396.90	9.08
503	21.0	396.90	5.64
504	21.0	393.45	6.48
505	21.0	396.90	7.88

[506 rows x 13 columns]

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2

..	...
501	22.4
502	20.6
503	23.9
504	22.0
505	11.9

Name: MEDV, Length: 506, dtype: float64

SPLITTING DATA INTO Training and Testing data

```
In [14]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [15]: print(X.shape, X_train.shape, X_test.shape)
```

```
(506, 13) (404, 13) (102, 13)
```

MODEL Training XGBoost Regressor

```
In [16]: # loading the model
model = XGBRegressor()
```

```
In [17]: # fitting is nothing but training of model
model.fit(X_train, Y_train)
```

```
Out[17]: XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_round
s=None,
             enable_categorical=False, eval_metric=None, feature_type
s=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=None, max_bin=None, max_cat_threshold=Non
e,
```

Evaluation

Prediction on training data

```
In [18]: # accuracy for prediction on training data
training_data_prediction = model.predict(X_train)
```

```
In [19]: print(training_data_prediction)
```

[23.124718	21.00754	20.102568	34.69276	13.904569	13.49714
21.997927	15.19248	10.902376	22.7026	13.800668	5.5908513
29.806072	50.005272	34.89682	20.596664	23.386295	19.18905
32.691494	19.63137	26.9884	8.40349	46.001217	21.7111
27.08755	19.365828	19.286129	24.817802	22.611925	31.707855
18.541298	8.704057	17.404493	23.701723	13.304713	10.520918
12.70769	24.98351	19.686928	14.899053	24.209797	24.994987
14.897052	17.01417	15.603933	12.6952915	24.52194	15.007025
49.999977	17.510012	21.203285	32.003624	15.595356	22.898546
19.32731	18.687641	23.30319	37.200005	30.095251	33.104855
20.992231	50.002266	13.401404	5.007679	16.5074	8.395711
28.68154	19.493786	20.596518	45.400917	39.804905	33.41812
19.840513	33.39644	25.271023	49.998375	12.517453	17.421158
18.604883	22.601322	50.00689	23.780687	23.312428	23.099342
41.69992	16.099009	31.596653	36.08056	6.999861	20.382881
20.000896	11.997474	24.995806	49.99009	37.89816	23.101585
41.27909	17.601559	16.308878	30.048021	22.85182	19.788818
17.106504	18.901857	18.939293	22.594484	23.15195	33.202766
14.999921	11.704477	18.80795	20.797966	17.995903	19.641762
50.00197	17.195467	16.402473	17.50968	14.601552	33.097614
14.495126	43.805664	34.900345	20.398077	14.625136	8.094603
11.778059	11.81811	18.70215	6.2982407	23.978172	13.066449
19.607876	49.999638	22.320127	18.919546	31.195917	20.702322
32.199364	36.165913	14.21831	15.7038765	49.983604	20.426397
16.174377	13.4097595	50.013714	31.606308	12.291674	19.217142
29.799698	31.501135	22.799213	10.192999	24.081253	23.703596
21.992378	13.785494	28.397635	33.19623	13.12277	19.048325
26.577816	36.955544	30.789625	22.790836	10.191355	22.198246
24.489767	36.18868	23.104813	20.114384	19.497326	10.80379
22.674639	19.507032	20.12285	9.60392	42.800453	48.79566
13.093994	20.288769	24.748339	14.103552	21.705412	22.232885
33.000233	21.116417	25.001814	19.119667	32.39981	13.608207
15.087661	23.093206	27.494678	19.37617	26.487434	27.501059
28.713585	21.22839	18.686394	26.723293	14.007895	21.704782
18.39718	43.114174	29.09505	20.300016	23.70995	18.291485
17.193474	18.319866	24.398975	26.397472	19.099882	13.307144
22.17664	22.189608	8.53707	18.89726	21.794674	19.346647
18.196587	7.5117607	22.395323	20.004005	14.399867	22.498732
28.510593	21.637846	13.804338	20.498287	21.900301	23.091843
50.000694	16.212135	30.307535	49.994564	17.797531	19.062796
10.398314	20.387184	16.501936	17.185501	16.736221	19.513783
30.499746	29.00148	19.556606	23.182762	24.39962	9.502657
23.892439	49.99509	21.199549	22.610487	19.990492	13.402443
19.960081	17.11145	12.723276	22.999899	15.244268	20.579155
26.206219	18.06881	24.096706	14.097596	21.700525	20.07267
25.013454	27.908047	22.916878	18.493921	22.197021	24.003748
14.810876	19.888756	24.399963	17.793882	24.5913	32.00859
17.793465	23.326887	16.09375	13.007798	10.997188	24.316023
15.592238	35.205746	19.599157	42.30329	8.796408	24.391514
14.101722	15.392906	17.298933	22.118727	23.100609	44.811054
17.803215	31.501413	22.80798	16.863302	23.907934	12.075618
38.7021	21.406189	16.001757	23.913858	11.895336	24.956305
7.200167	24.696964	18.20617	22.466045	23.026604	24.300848
17.110569	17.80117	13.496556	27.07502	13.305881	21.898989

19.997826	15.361549	16.58387	22.302023	24.720388	21.399815
22.89386	29.597786	21.871275	19.895018	29.599123	23.39317
13.804395	24.44912	11.906414	7.2283416	20.504478	9.699921
48.29024	25.197203	11.686412	17.400322	14.49293	28.59803
19.38137	22.44903	7.014273	20.590538	22.981966	19.69853
23.695066	25.00565	28.003778	13.382997	14.525479	20.311686
19.308327	24.096216	14.900837	26.39965	33.302242	23.6273
24.596827	18.503511	20.899784	10.401823	23.296179	13.112332
24.671389	22.596157	20.502825	16.81155	10.204616	33.805656
18.616283	49.996334	23.796791	23.900902	21.182632	18.816933
8.505743	21.498335	23.196657	21.019245	16.606592	28.09294
21.20949	28.39348	14.284869	50.002743	30.989561	24.997051
21.427883	19.001093	29.003275	15.20375	22.797543	21.770397
19.91061	23.77714]			

```
In [20]: # for knowing the error use R squared error
# i just calculate the difference btwn the value predicted by model and origin
score_1 = metrics.r2_score(Y_train, training_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_train, training_data_prediction)

print("R squared error : ", score_1) # this value shuld be near to 1 if it is
# if 0 then model preforming perfectly
print("Mean Absolute Error : ", score_2)
```

R squared error : 0.9999970846673867
Mean Absolute Error : 0.010642476601175762

```
In [21]: # accuracy for prediction on testing data
test_data_prediction = model.predict(X_test)
```

visualizing the actual and predicted prices

```
In [25]: plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Prices for teste data")
plt.ylabel("Predicted Prices test data")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```



```
In [22]: plt.scatter(Y_test, test_data_prediction)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```



Finding the error value for test data

```
In [ ]: # for knowing the error use R squared error
# i just calculate the difference btwn the value predicted by model and origin
score_1 = metrics.r2_score(Y_test, test_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_test, test_data_prediction)

print("R squared error : ", score_1) # this value shuld be near to 1 if it is
# if 0 then model preforming perfectly
print("Mean Absolute Error : ", score_2)
```

R squared error : 0.8950119616414808
Mean Absolute Error : 2.2481449632083668